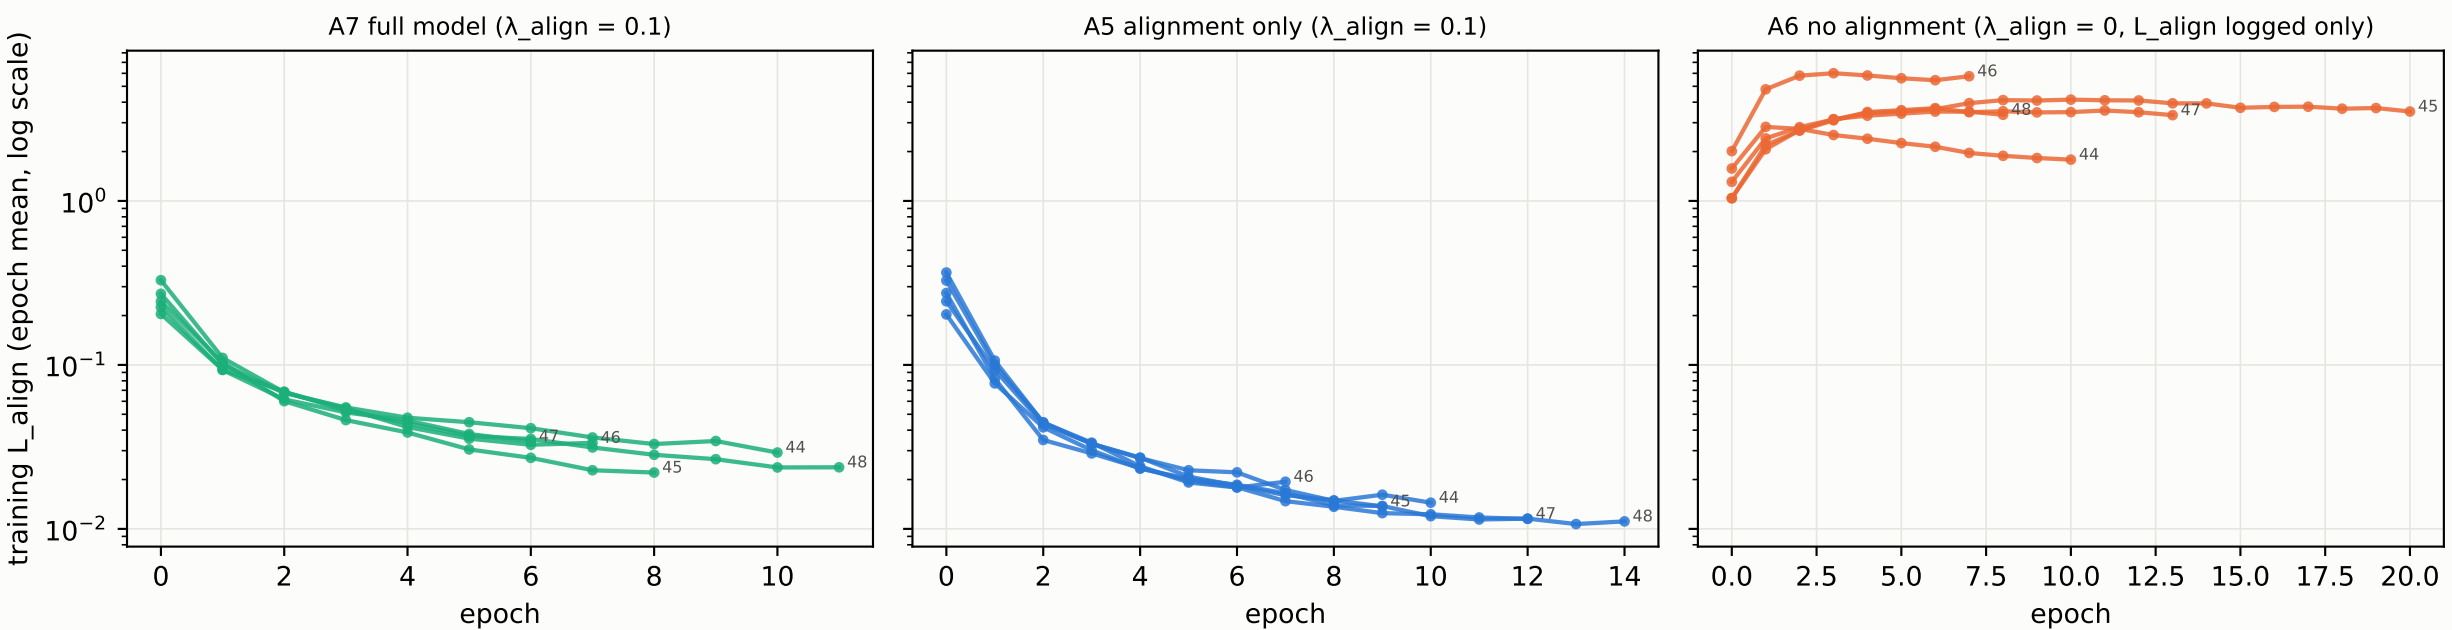
Alignment loss per epoch, every seed (44–48), canonical 2026-10-04 sweep; labels = seed
A7 full model (λ_align = 0.1)
A5 alignment only (λ_align = 0.1)
A6 no alignment (λ_align = 0, L_align logged only)
training L_align (epoch mean, log scale)
epoch
epoch
epoch